

Same 20-seed ensembles, same test events: opposite verdicts on the two axes (up = pre-training better; bars: paired-bootstrap $\oplus$ seed/ $\sqrt{K}$)

